

The used car market has significantly grown in recent times, with clients ranging from used car dealers and buyers. You are provided with a car evaluation dataset that has features like price, doors, safety, and so on.

Objective: You are required to create a robust model that allows stakeholders to predict the condition of a used vehicle.

Actions to Perform:

Predict the condition of a vehicle based on its features. Plot the most important features. Train multiple classifiers and compare the accuracy. Evaluate the XGBoost model with K-fold cross-validation.

It is a multi-class classification problem. Task is to classify our dataset into 4 classes : Unacceptable, Acceptable, Good, Very-Good.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: data=pd.read_csv('car_evaluation.csv')
```

```
In [3]: data.head()
```

```
Out[3]:
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

Independent variables

buying : buying price

maint : maintenance price

doors : number of doors

persons : capacity in terms of persons to carry

lug_boot : the size of luggage boot

safety: estimated safety of the car

Target Variable

Class

In [5]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   object
1   maint       1728 non-null   object
2   doors       1728 non-null   object
3   persons     1728 non-null   object
4   lug_boot    1728 non-null   object
5   safety      1728 non-null   object
6   class       1728 non-null   object
dtypes: object(7)
memory usage: 94.6+ KB
```

Columns are categorical. We are looking for unique values on each column

In [6]: `for i in data.columns:`
`print(data[i].unique(), "\t", data[i].nunique())`

```
['vhigh' 'high' 'med' 'low']      4
['vhigh' 'high' 'med' 'low']      4
['2' '3' '4' '5more']            4
['2' '4' 'more']                  3
['small' 'med' 'big']             3
['low' 'med' 'high']              3
['unacc' 'acc' 'vgood' 'good']    4
```

Checking the distribution of these unique categories among the columns

```
In [7]: for i in data.columns:
        print(data[i].value_counts())
        print()
```

```
med      432
low      432
vhigh    432
high     432
Name: buying, dtype: int64
```

```
med      432
low      432
vhigh    432
high     432
Name: maint, dtype: int64
```

```
4         432
5more     432
2         432
3         432
Name: doors, dtype: int64
```

```
4         576
more      576
2         576
Name: persons, dtype: int64
```

```
big       576
med       576
small     576
Name: lug_boot, dtype: int64
```

```
med       576
low       576
high      576
Name: safety, dtype: int64
```

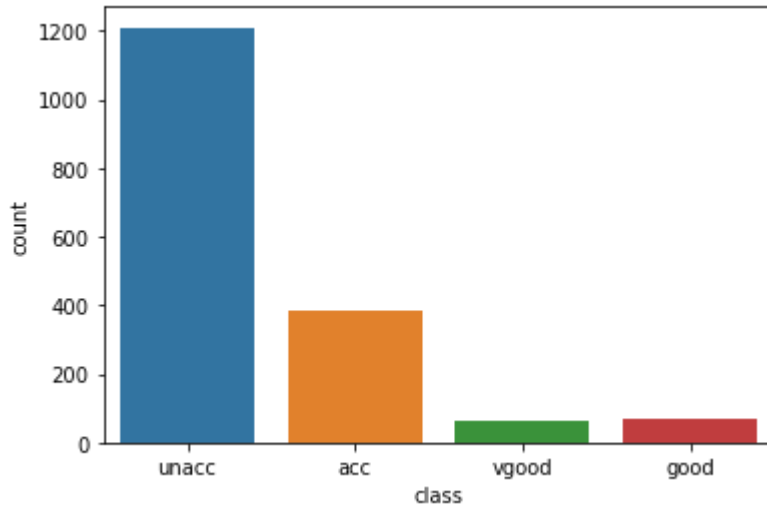
```
unacc    1210
acc       384
good       69
vgood     65
Name: class, dtype: int64
```

From the above we can see that all columns except "class" are distributed equally among data

Class Distribution

```
In [8]: sns.countplot(data['class'])
```

```
Out[8]: <AxesSubplot:xlabel='class', ylabel='count'>
```



From above we see that 'class' is unbalanced with larger values under 'unacc'. So, this is an unbalanced multiclass classification problem.

Dummy Encoding

```
In [9]: from sklearn.preprocessing import LabelEncoder
```

```
In [10]: le=LabelEncoder()
```

```
In [11]: for i in data.columns:  
         data[i]=le.fit_transform(data[i])
```

```
In [12]: data.head()
```

```
Out[12]:
```

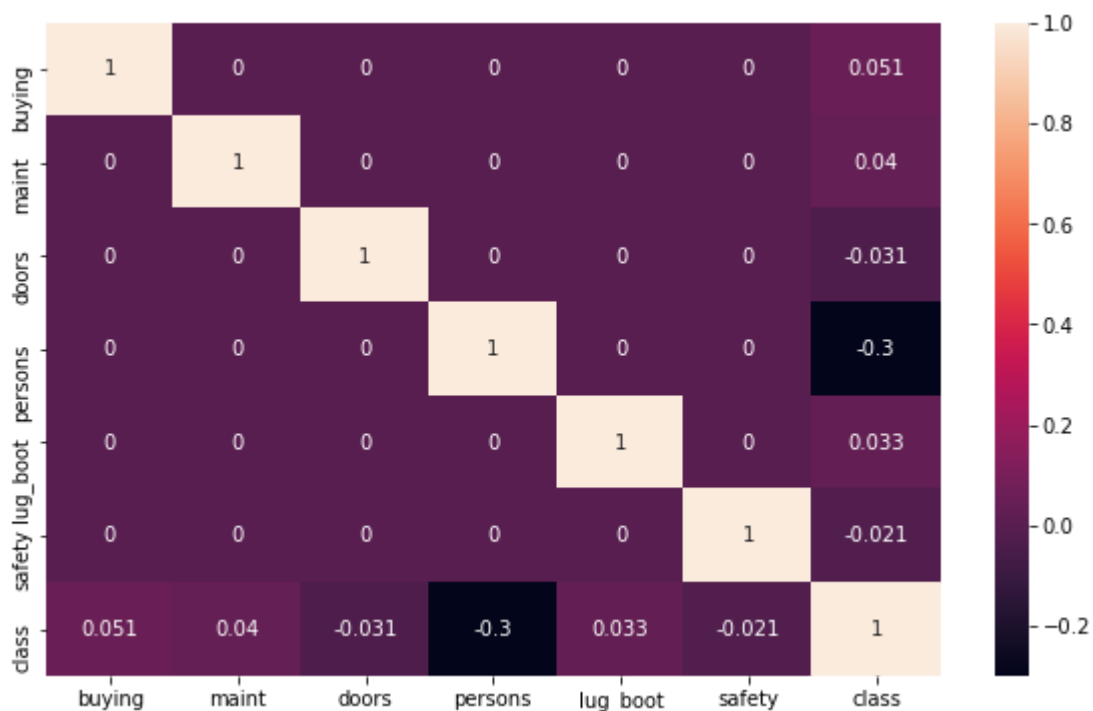
	buying	maint	doors	persons	lug_boot	safety	class
0	3	3	0	0	2	1	2
1	3	3	0	0	2	2	2
2	3	3	0	0	2	0	2

	buying	maint	doors	persons	lug_boot	safety	class
3	3	3	0	0	1	1	2
4	3	3	0	0	1	2	2

Correlation Matrix

```
In [13]: fig=plt.figure(figsize=(10,6))
sns.heatmap(data.corr(),annot=True)
```

```
Out[13]: <AxesSubplot:>
```



Most of columns have a very weak correlation with "class" so, doing any analysis on them may not give a productive output

```
In [14]: X=data[data.columns[:-1]]
Y=data['class']
```

```
In [15]: X.head()
```

Out[15]:

	buying	maint	doors	persons	lug_boot	safety
0	3	3	0	0	2	1
1	3	3	0	0	2	2
2	3	3	0	0	2	0
3	3	3	0	0	1	1
4	3	3	0	0	1	2

```
In [16]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.3, random_state=42)
```

Model Selection

```
In [18]: from sklearn.model_selection import learning_curve
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
```

1. Logistic Regression

```
In [19]: logreg=LogisticRegression(solver='newton-cg',multi_class='multinomial')
```

```
In [20]: logreg.fit(X_train,Y_train)
```

```
Out[20]: LogisticRegression(multi_class='multinomial', solver='newton-cg')
```

```
In [21]: pred=logreg.predict(X_test)
```

```
In [22]: logreg.score(X_test,Y_test)
```

```
Out[22]: 0.6647398843930635
```

Logistic Regression is giving a very low accuracy. Will check with other algorithms

2. KNN Classifier

```
In [23]: knn=KNeighborsClassifier(n_jobs=-1)
```

```
In [24]: knn.fit(X_train,Y_train)
pred=knn.predict(X_test)
knn.score(X_test,Y_test)
```

```
Out[24]: 0.9017341040462428
```

```
In [26]: print(classification_report(Y_test,pred))
```

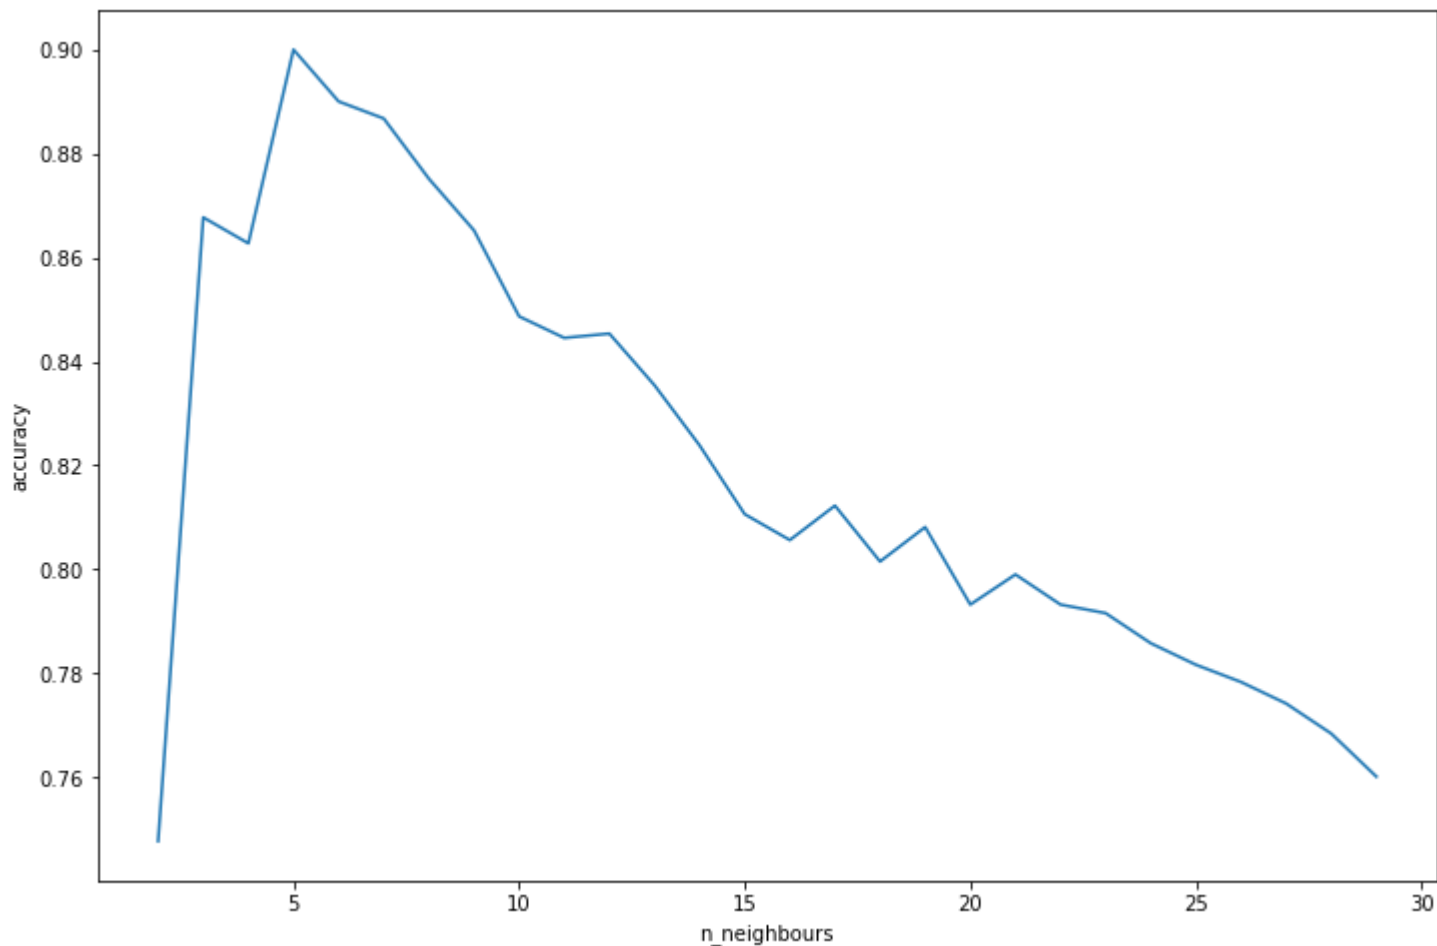
	precision	recall	f1-score	support
0	0.82	0.79	0.80	118
1	0.77	0.53	0.62	19
2	0.93	0.99	0.96	358
3	1.00	0.50	0.67	24
accuracy			0.90	519
macro avg	0.88	0.70	0.76	519
weighted avg	0.90	0.90	0.90	519

Accuracy isn't a good criterion to evaluate unbalanced classification, so check 'f1-score' f1-score is 0.9 which is better than previous model

```
In [27]: avg_score=[]
for k in range(2,30):
    knn=KNeighborsClassifier(n_jobs=-1,n_neighbors=k)
    score=cross_val_score(knn,X_train,Y_train,cv=5,n_jobs=-1,scoring='accuracy')
    avg_score.append(score.mean())
```

```
In [29]: plt.figure(figsize=(12,8))
plt.plot(range(2,30),avg_score)
plt.xlabel("n_neighbours")
plt.ylabel("accuracy")
#plt.xticks(range(2,30,2))
```

```
Out[29]: Text(0, 0.5, 'accuracy')
```



n_neighbours=5 is giving better accuracy as well as f1-score for our data So, with KNN Classification we achieve accuracy of 90%

3. Random Forests Classifier

```
In [30]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import f1_score
```

```
In [33]: rfc=RandomForestClassifier(n_jobs=-1,random_state=51)
```

```
In [36]: rfc.fit(X_train,Y_train)  
print(rfc.score(X_test,Y_test))  
print(f1_score(Y_test,rfc.predict(X_test),average='macro'))
```

0.9730250481695568

0.9245337130459484

RFC gives 95% accuracy

In []: